

High Volume Ratelimiting with Redis

by

Triton Ho



Content

- Our problem
- v0 approach
- v1 approach

Our problem

- Unpredictable events cause sudden burst of QPS
- Auto-scaling is not fast enough
- Database is painful to scale
 - Especially during peak hours
 - Especially with m***odb
- Solution: HTTP 429, Too Many Requests

Why not Nginx

- Users play 30mins game in our app
 - Interruption causes horribly bad UX
- When QPS over server capacity
 - Allow requests from existing users
 - Block requests from new users
- new users = no request in previous 30mins
- Nginx doesn't support such rate-limiting

Technical Requirements

- When total active user $> X$, system enters "throttling mode"
- During throttling mode
 - Check for AccessToken
 - If the AccessToken is used in previous 30mins
Allows it
else
HTTPS 429

V0 approach in Redis

- hs_20190309:2300
 - Init TTL: 60min
 - KeyValues
 - TokenA: 20190309:231012
 - TokenB: 20190309:231839
 - TokenC: 20190309:232912
- hs_20190309:2330
 - Init TTL: 60min
 - KeyValues
 - TokenB: 20190309:233139
 - TokenC: 20190309:233312
 - TokenD: 20190309:233454

v0 approach

- Simple, but not working
- Maintain multiple hashset in redis
 - Key = AccessToken, Value = lastUsedTs
- Updates lastUsedTs in every request
- Active User \sim $\text{math.Max}(\text{HLEN}(\text{hs_current}), \text{HLEN}(\text{hs_previous}))$
- During system throttling, check for “current” and “previous” hashset

Sound bad, and epic fails in v0

- Rate-limiting's goal
 - protecting the redis(and other components) from QPS burst
- But in v0 design
 - ALL requests need one redis call before rejected
- All requests need update on single hashSet
 - Super hotspot in redis cluster

Goal in v1

- The redis call MUST be independent of QPS
- Use local memory in application server instead of redis
- **Constant** local memory usage
- Seems mission-impossible.....

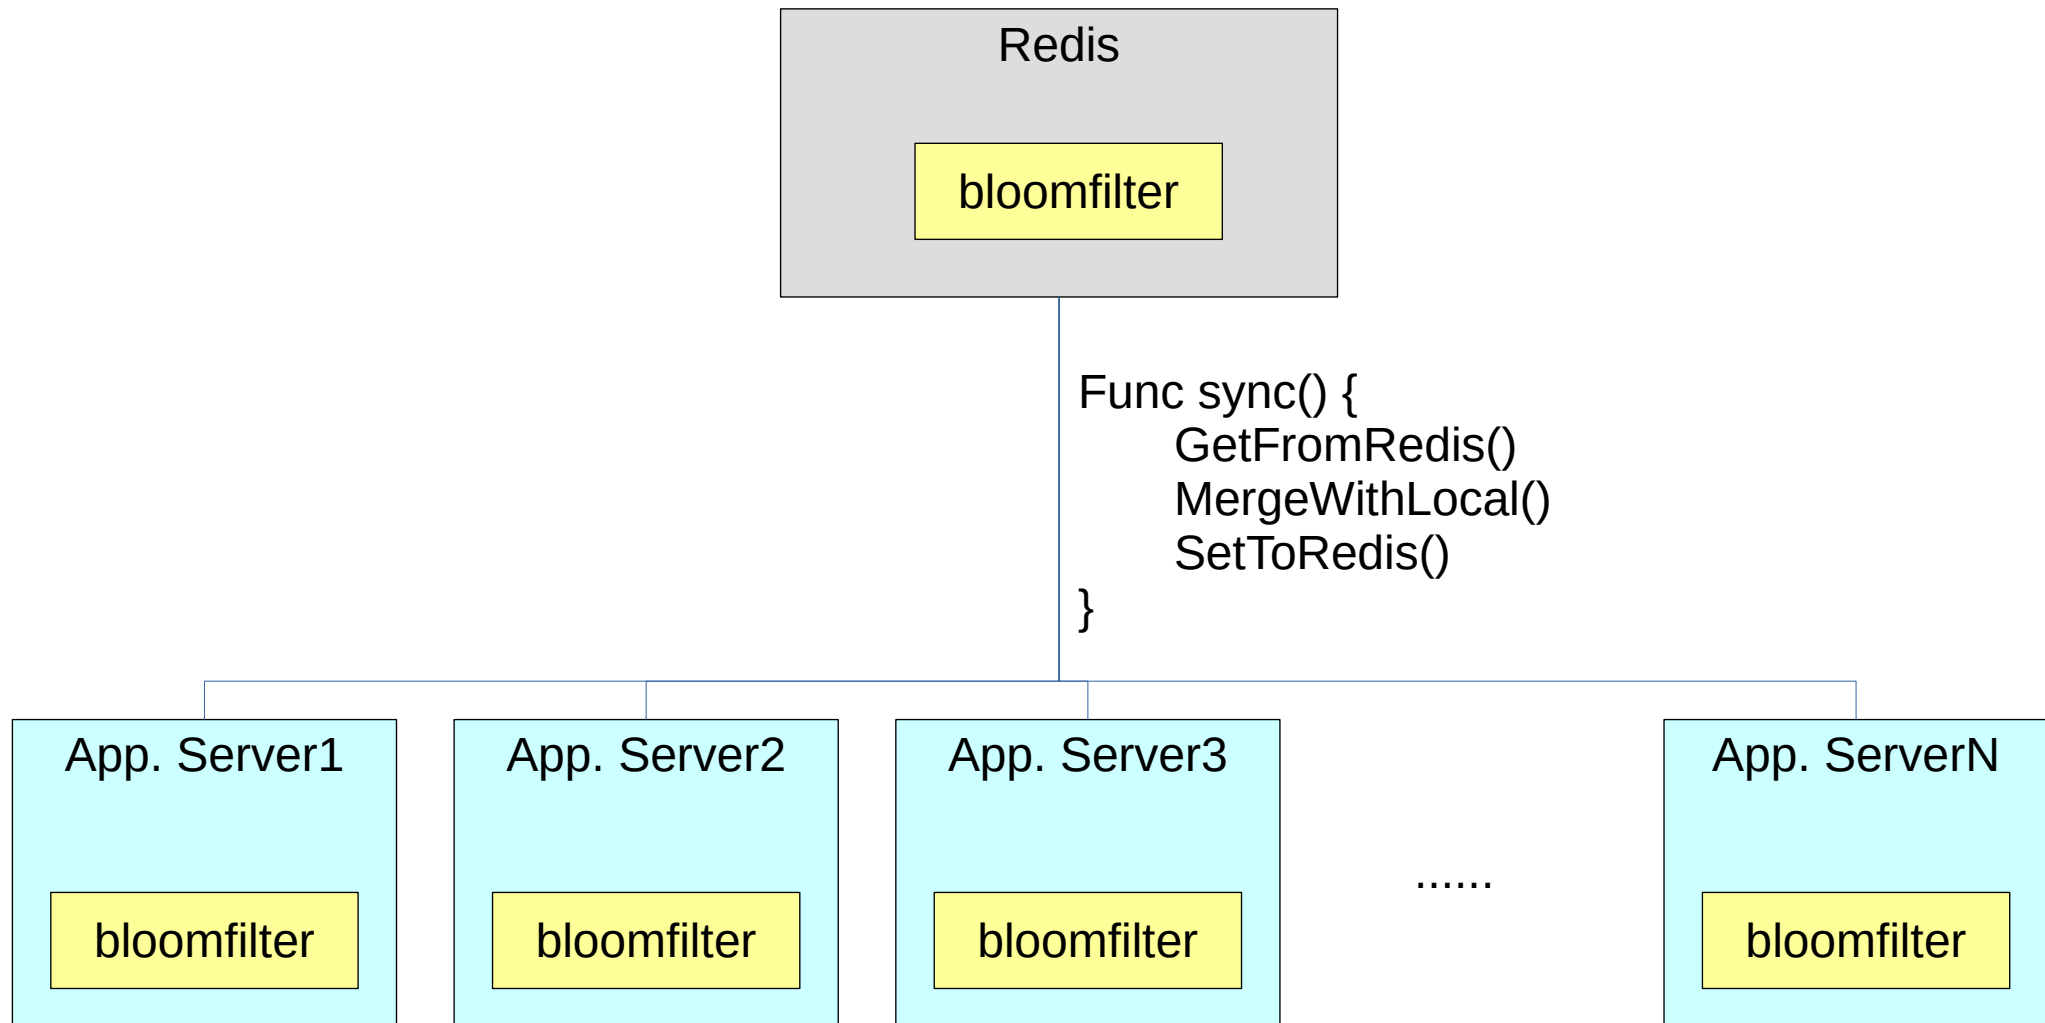
Bloomfilter data structure

- Constant memory usage
- Add(string s)
 - Add string s into the bloomfilter
- Test(string s) bool
 - Test if string s exists
 - may return false positive
- AppxCount() int
 - Return the approximation of the number of stored items
- Merge(Bloomfilter bf)
 - Merge two bloomfilter

Why bloomfilter helpful?

- Little false positive doesn't matter
- AppxCount() gives fast approximation of active users in the system
- Merge use bitwise OR operation on byte array
 - Allow super fast synchronization between application server
 - Gzipped byte array is small, small network bandwidth requirement

V1 (simplified) architecture



App. server Synchronization

- Use redis as central storage
 - No extra backend component
a.k.a. no extra approval from manager, smile~
 - Every second, the app. server sync the local copy with redis
- Optimistic lock is used to protect concurrent sync from multiple app. server
 - HashSet is used instead of KeyValue in redis
 - “ts” field for timestamp,
 - “value” field for bloomfilter content
 - Evalsha() is needed to run the lua script in redis

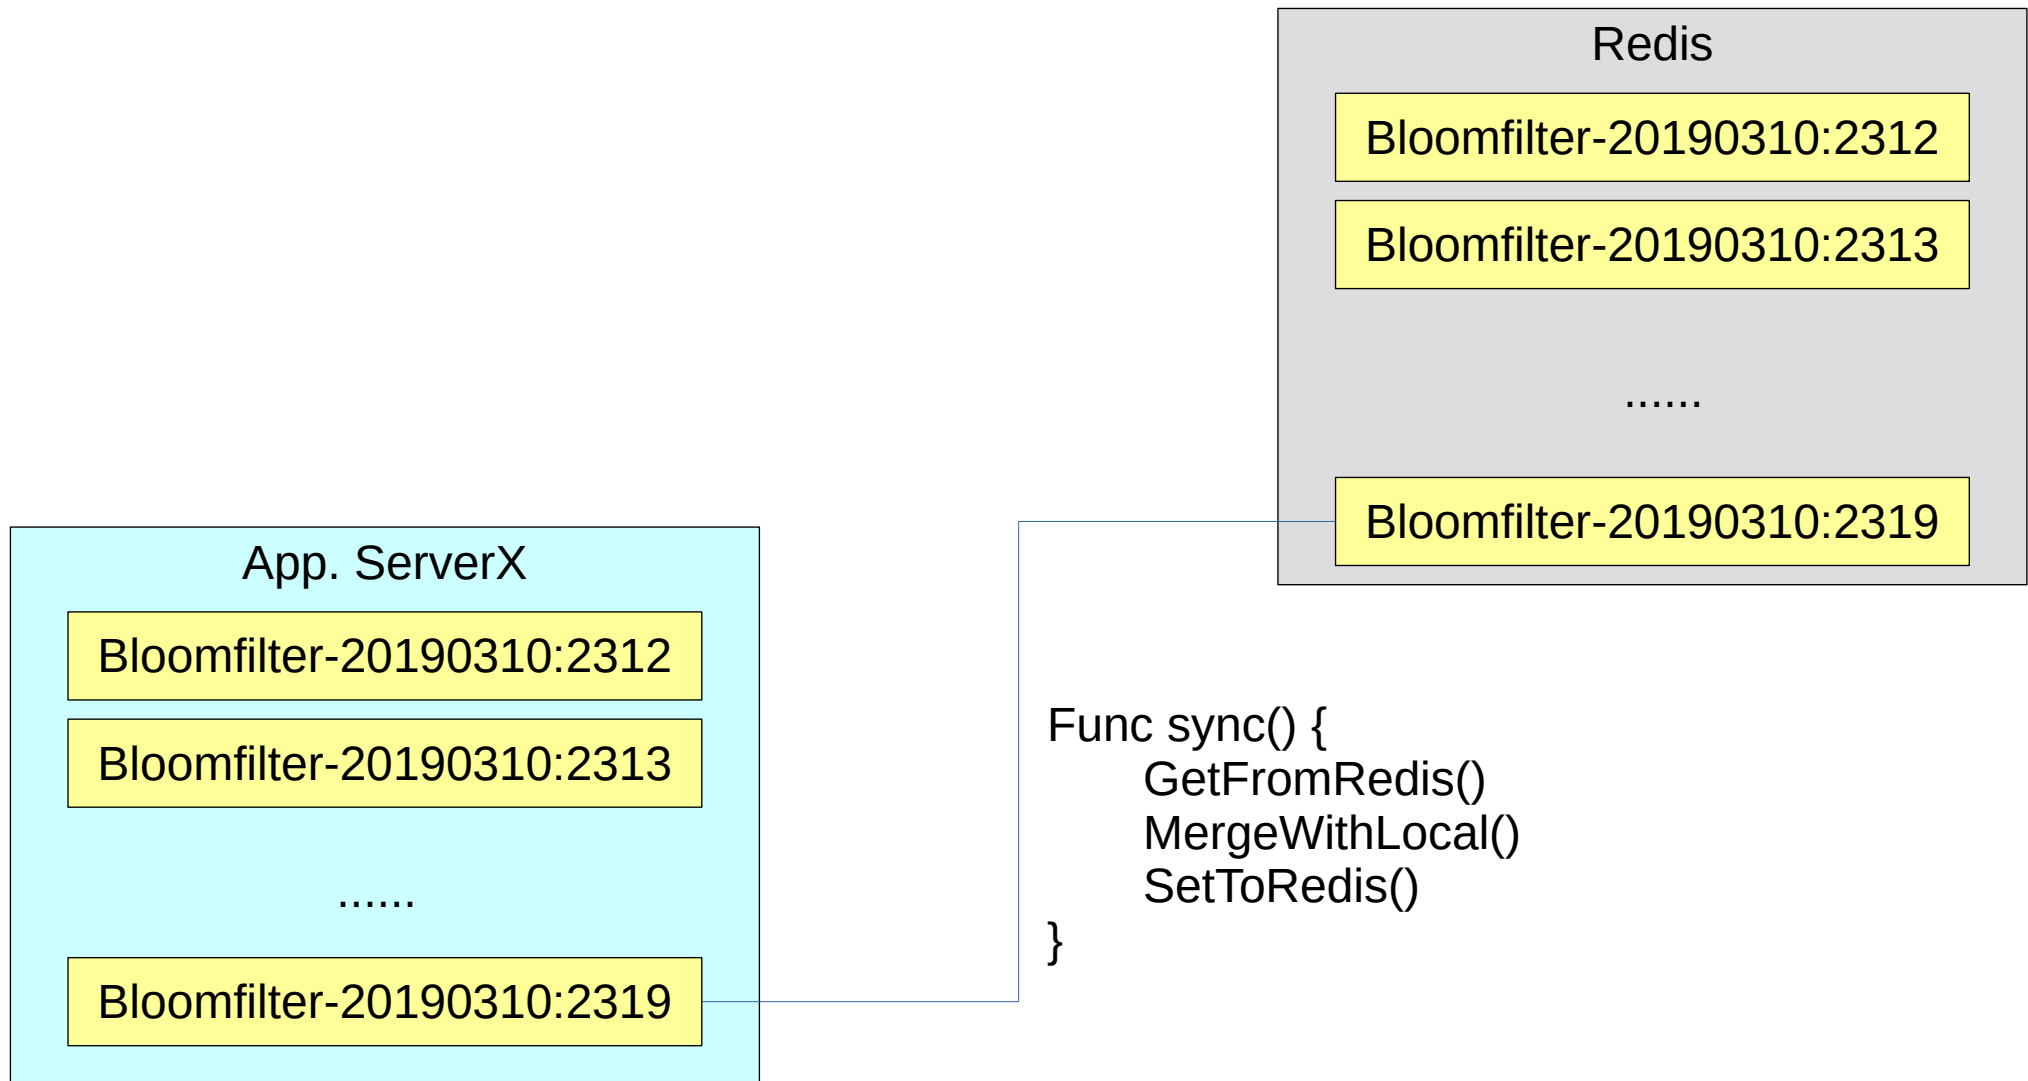
Concurrency in app. server

- RW-lock is used
 - Allow concurrency read, but single-thread update of bloomfilter
- To minimize latency, the update of bloomfilter is deferred
 - Buffered channel in Golang is great~

Removing old AccessToken

- Bloomfilter has no “expiration” mechanism
- “log rotation” on bloomfilter
 - Keep 30 bloomfilter
 - in every minute, remove the oldest one and create a new one
- New AccessToken add to the latest bloomfilter
 - Thus, only latest bloomfilter need to keep sync. with redis

V1 architecture



fails in v1

- Optimistic lock description in wikipedia
 - OCC is generally used in environments with low data contention
 - if contention for data resources is frequent, the cost of repeatedly restarting transactions hurts performance significantly
- The sync. conflict increase exponential with app. server number.
 - Caused abnormal high bandwidth usage during peak hours

v1.1 modification

- During optimistic lock conflict, the app. server simply give up the SetToRedis()
 - The sync. is done in every second, some failed SetToRedis() can be tolerated

Lessons learnt

- Performing redis call per request for ratelimiting is horribly bad idea
 - Your redis will fail during DDoS
 - Similarly, random-string AccessToken is BAD
- Be careful for hotspot in redis cluster
- Optimistic lock is suitable if the conflict rate is very low
- Algorithm and data structure has practical use
 - DO NOT SLEEP IN UNIVERSITY COURSES

More thought

- In the world of >10K QPS, consistency is usually sacrificed for performance
 - Bloomfilter is an appx algorithm
- Efficient use of app. server local memory is the key for high performance

Extras

- Can you use redis to implement a priority task queue?
 - Sorted Set should be avoided at all cost
 - For fault tolerance application, you should use rabbitMQ instead

End~

